

ClickHouse as a High-Performance Monitoring System, from zero to hero

DevOpsAI SG 2025



Welcome!

Agenda

01

Introduction

02

Pre-Requisites

03

ClickHouse Fundamentals

04

Data Modeling - basics

05

Python - data generator

06

Visualization - Grafana

07

Python ML

08

FAQ

09

10

11

12



STAND A CHANCE TO WIN AN R2-D2 LEGO SET!

Scan the **QR code**
and put your
name in for the
lucky draw!

Meet us at the
ClickHouse-sponsored coffee
cart during afternoon tea break!





Introduction

Who am I ???



||||· ClickHouse

The definition of observability

Definition (software)

"In software engineering, more specifically in distributed computing, observability is the **ability to collect data** about programs' execution, modules' internal states, and the communication among components. To improve observability, software engineers use a wide range of **logging and tracing techniques** to gather telemetry information, and **tools to analyze and use it**. Observability is **foundational** to **site reliability engineering**, as it is the first step in **triaging a service outage**. One of the goals of observability is to minimize the amount of prior knowledge needed to debug an issue"

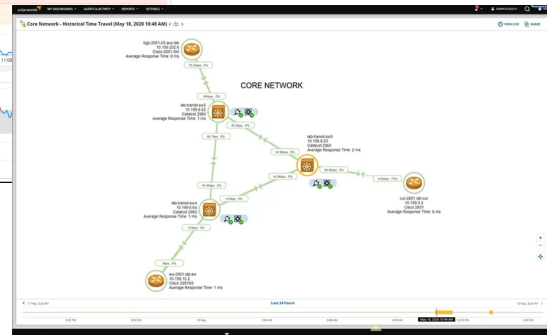

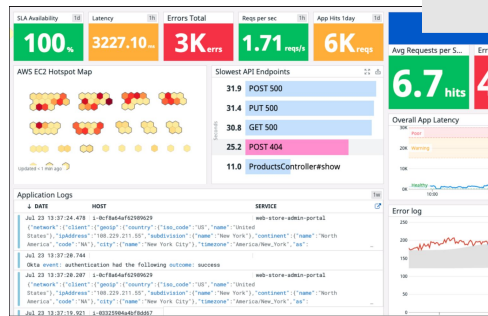
Quoted from [https://en.wikipedia.org/wiki/Observability_\(software\)](https://en.wikipedia.org/wiki/Observability_(software))



What is happening in the field?

- There are many observability platforms (data dog, splunk, dyntrace, sumologic), aren't they enough for covering everything?

- **Detour:**
 - costs,
 - transparency,
 - speed (performance),
 - complexity vs flexibility



Yet another product / platform in the field...

ClickHouse - new kid on the block

- Costs - we might not be the cheapest but definitely not expensive ([price calculator](#))
- Transparency - our [pricing](#) is transparent, the data stored is totally visible in table format (SQL), no middleware in between, can fit into workflows through [connectors](#)
- Performance - we are fast ([ClickBench](#)) in both ingestion and retrieval; storage is efficient (lz4 by default, also available for [zstd](#))
- Complexity - check how Splunk and others' UI were, ClickHouse is simply SQL and tables which are developer friendly (previous slide)



What is ClickHouse actually?

Open source **column-oriented** **distributed** **OLAP database**

Since 2009
35,000+ GitHub stars
1300+ contributors
500+ releases

Best for aggregations
Files per column
Sorting and indexing
Background merges

Replication
Sharding
Multi-master
Cross-region

Analytics use cases
Aggregations
Visualization
Mostly immutable data



Objective of the workshop

Illustrate what a typical (simplified) observability workflow looks like and how ClickHouse fits into the flow.

Throughout the journey:

- Understand the basics of ClickHouse (e.g. Schema, Materialized Views)
- Extensibility through integrations (e.g. with Python client and Grafana)



02

Pre Requisites

Tools before you start (15 min)

Instructions (for Mac only) available at <https://shorturl.at/BO0HO>

- Docker Desktop
- Docker Compose (for older version of Docker Desktop that does not come with `compose`)
- (optional) Git
- (optional) Python3
- A modern web browser (chrome, firefox)

And...

We are all GOOD !





ClickHouse fundamentals

Understanding Columnar Storage

- One reason ClickHouse is so fast is because it is *column-oriented*
- But you need to understand columnar storage to get the benefit

So what does it mean?



Concepts

- In ClickHouse, per table required a table engine; typically we will be using [MergeTree](#)(s) which is optimized for high performance inserts and queries (should be good enough for most cases)

Primary key

- Primary key is different...

The primary key determines how the data is stored and searched

- **ORDER BY** used if no **PRIMARY KEY**
- *Not unique to each row*

You can also define the primary key using **ORDER BY**

```
CREATE TABLE my_table
(
    column1    FixedString(1),
    column2    UInt32,
    column3    String
)
ENGINE = MergeTree()
PRIMARY KEY (column1, column2)
```



Concepts (cont')

- Data are retrieved in Granule(s)
 - **granule** – a logical breakdown of rows inside an uncompressed block; default is 8,192 rows
 - **primary key** – the sort order of a table
 - **primary index** – an in-memory index containing the values of the primary keys of the first row of each granule
 - **part** – a folder of files consisting of the column files and index file of a subset of a table's data

A granule is the smallest indivisible amount of data that ClickHouse reads when searching rows





Data Modeling - Basics

Schema

- In ClickHouse, table Schema (aka DDL) is important.
- It is a big-data oriented SQL db (not **noSQL** db) check it out on Elasticsearch too!
- It is possible to alter the existing columns but have to pay a price; it is done through `mutation`
- Instead of updating the schema, we could consider Materialized views



Materialized View

Use cases:


- If your source / raw data required to be processed (e.g. extraction of certain fields from a json data structure)
- If the existing table's Index (sorting-key / primary-key) does not meet a use case's requirements
- To filter out data from the source table into dedicated tables (e.g. forward a data row based on the `region` field's value, each region would have its own table)
- To consolidate various source tables' data into a single target table



Materialized View

- General link : <https://shorturl.at/o94Oy>
- Refreshable MV : <https://shorturl.at/SbbPw>
- Incremental MV : <https://shorturl.at/lxJRy>





FAQ

Who am I ???



||||· ClickHouse

STAND A CHANCE TO WIN AN R2-D2 LEGO SET!

Scan the **QR code**
and put your
name in for the
lucky draw!

Meet us at the
ClickHouse-sponsored coffee
cart during afternoon tea break!



Thank you!